

Apache Phoenix (SQL on HBase) Installation

Tao Wu, 06/18/2015

1 What is Apache Phoenix

Apache Phoenix is a relational database layer over HBase delivered as a client-embedded JDBC driver targeting low latency queries over HBase data. Apache Phoenix takes your SQL query, compiles it into a series of HBase scans, and orchestrates the running of those scans to produce regular JDBC result sets. The table metadata is stored in an HBase table and versioned, such that snapshot queries over prior versions will automatically use the correct schema. Direct use of the HBase API, along with coprocessors and custom filters, results in **performance** on the order of milliseconds for small queries, or seconds for tens of millions of rows.

<https://phoenix.apache.org/>

2 Who is using Apache Phoenix

It is open-sourced by Salesforce.com, now is top-level apache project, and Cloudera announced their support on Apache Phoenix in their labs. Hortonworks also integrates Phoenix in HDP.

https://phoenix.apache.org/who_is_using.html

<http://blog.cloudera.com/blog/2015/05/apache-phoenix-joins-cloudera-labs/>

Who is using Apache Phoenix? Read more [here...](#)



3 Installation and Short Try

3.1 Prepare package.

3.1.1 Checkout the Cloudera version of Phoenix

<https://github.com/cloudera-labs/phoenix/tree/master>

Checkout the commit `c870bd8e1cb27d7d8e2bc954a6ff1b528a9177bd` for version `CLOUDERA-BUILD` Fixing the incompats against CDH5.4 HBase

```
git checkout c870bd8e1cb27d7d8e2bc954a6ff1b528a9177bd
```

3.1.2 Maven Build

```
mvn package -Dskiptests
```

Fix the maven build problem.

```
[ERROR] /home/taowu/git-newegg/phoenix/phoenix-core/src/it/java/org/apache/hadoop/hbase/regionserver/wal/WALReplayWithIndexWritesAndCompressedWALIT.java:[101,49] cannot access org.apache.hadoop.hdfs.DistributedFileSystem class file for org.apache.hadoop.hdfs.DistributedFileSystem not found
```

add Hadoop-hdfs dependency in pom.xml to fix this problem

```
+ <dependency>
+   <groupId>org.apache.hadoop</groupId>
+   <artifactId>hadoop-hdfs</artifactId>
+   <version>${hadoop-two.version}</version>
+ </dependency>
```

3.2 Deployment

3.2.1 Deploy phoenix jar to region servers.

Copy these 2 jars to hbase lib folder

```
phoenix-4.3.0-cdh5.4.0-SNAPSHOT-server.jar
```

```
phoenix-core-4.3.0-cdh5.4.0-SNAPSHOT.jar
```

The HBase lib folder should be like this for Cloudera Manager Installation:

```
/opt/cloudera/parcels/CDH-5.4.0-1.cdh5.4.0.p0.27/lib/hbase/lib
```

```
sudo chmod 777 phoenix-*.jar
```

3.2.2 Restart HBase cluster.

3.3 Command Line Client Deployment

Note: Should run this client by JDK 7

untar the phoenix-4.3.0-cdh5.4.0.tar.gz

3.3.1 Try sqlline.py to verify the installation.

```
./sqlline.py server11,server03,server01,server04,server02
```

```
[tw79@e3ecmrhdp01 bin]$ ./sqlline.py e3ecmrhdp11.mercury.corp,e3ecmrhdp03.mercury.corp,e3ecmrhdp01.mercury.corp,e3ecmrhdp04.mercury.corp,e3ecmrhdp02.mercury.corp
Setting property: [isolation, TRANSACTION_READ_COMMITTED]
Issuing: !connect jdbc:phoenix:e3ecmrhdp11.mercury.corp,e3ecmrhdp03.mercury.corp,e3ecmrhdp01.mercury.corp,e3ecmrhdp04.mercury.corp,e3ecmrhdp02.mercury.corp none non
s org.apache.phoenix.jdbc.PhoenixDriver
Connecting to jdbc:phoenix:e3ecmrhdp11.mercury.corp,e3ecmrhdp03.mercury.corp,e3ecmrhdp01.mercury.corp,e3ecmrhdp04.mercury.corp,e3ecmrhdp02.mercury.corp
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
15/06/18 12:05:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Connected to: Phoenix (version 4.3)
Driver: PhoenixEmbeddedDriver (version 4.3)
Autocommit status: true
Transaction isolation: TRANSACTION_READ_COMMITTED
Building list of tables and columns for tab-completion (set fastconnect to true to skip)...
155/155 (100%) Done
Done
sqlline version 1.1.8
0: jdbc:phoenix:e3ecmrhdp11.mercury.corp,e3ec> █
```

3.3.2 Try psql.py to import data to do test.

```
./psql.py server01,server02,server03,server04,server11 ../examples/WEB_STAT.sql
```

```
../examples/WEB_STAT.csv ../examples/WEB_STAT_QUERIES.sql
```

```
sqlline version 1.1.8
0: jdbc:phoenix:e3ecmrhdp11.mercury.corp,e3ec> select * from WEB_STAT;
+-----+-----+-----+-----+-----+
| HOST | DOMAIN | FEATURE | DATE | CORE |
+-----+-----+-----+-----+-----+
| EU | Apple.com | Mac | 2013-01-01 01:01:01.000 | 35 |
| EU | Apple.com | Store | 2013-01-03 01:01:01.000 | 345 |
| EU | Google.com | Analytics | 2013-01-13 08:06:01.000 | 25 |
| EU | Google.com | Search | 2013-01-09 01:01:01.000 | 395 |
| EU | Salesforce.com | Dashboard | 2013-01-06 05:04:05.000 | 12 |
| EU | Salesforce.com | Login | 2013-01-12 01:01:01.000 | 5 |
| EU | Salesforce.com | Reports | 2013-01-02 12:02:01.000 | 25 |
| EU | Salesforce.com | Reports | 2013-01-02 14:32:01.000 | 125 |
| EU | Salesforce.com | Reports | 2013-01-05 03:11:12.000 | 75 |
| EU | Salesforce.com | Reports | 2013-01-05 04:14:12.000 | 475 |
| EU | Salesforce.com | Reports | 2013-01-13 08:04:04.000 | 355 |
| NA | Apple.com | Login | 2013-01-01 01:01:01.000 | 35 |
| NA | Apple.com | Login | 2013-01-04 01:01:01.000 | 135 |
| NA | Apple.com | Mac | 2013-01-02 04:01:01.000 | 345 |
| NA | Apple.com | Mac | 2013-01-08 01:01:01.000 | 3 |
| NA | Apple.com | iPad | 2013-01-05 01:01:01.000 | 85 |
| NA | Apple.com | iPad | 2013-01-06 01:01:01.000 | 35 |
| NA | Apple.com | iPad | 2013-01-07 01:01:01.000 | 9 |
| NA | Google.com | Analytics | 2013-01-07 06:01:01.000 | 23 |
| NA | Google.com | Analytics | 2013-01-11 01:02:01.000 | 7 |
| NA | Google.com | Analytics | 2013-01-14 01:01:01.000 | 65 |
| NA | Google.com | Search | 2013-01-08 08:01:01.000 | 345 |
| NA | Google.com | Search | 2013-01-10 01:05:01.000 | 835 |
| NA | Google.com | Search | 2013-01-12 01:01:01.000 | 8 |
| NA | Salesforce.com | Dashboard | 2013-01-03 11:01:01.000 | 88 |
| NA | Salesforce.com | Dashboard | 2013-01-11 01:01:01.000 | 335 |
| NA | Salesforce.com | Dashboard | 2013-01-14 04:07:01.000 | 5 |
| NA | Salesforce.com | Login | 2013-01-01 01:01:01.000 | 35 |
| NA | Salesforce.com | Login | 2013-01-04 06:01:21.000 | 3 |
| NA | Salesforce.com | Login | 2013-01-04 11:01:11.000 | 23 |
| NA | Salesforce.com | Login | 2013-01-08 14:11:01.000 | 345 |
| NA | Salesforce.com | Login | 2013-01-10 01:01:01.000 | 345 |
| NA | Salesforce.com | Login | 2013-01-16 01:01:01.000 | 785 |
| NA | Salesforce.com | Login | 2013-01-17 01:01:01.000 | 355 |
| NA | Salesforce.com | Login | 2013-01-17 02:20:01.000 | 1235 |
| NA | Salesforce.com | Reports | 2013-01-09 16:33:01.000 | 35 |
| NA | Salesforce.com | Reports | 2013-01-09 17:36:01.000 | 355 |
| NA | Salesforce.com | Reports | 2013-01-15 04:09:01.000 | 65 |
| NA | Salesforce.com | Reports | 2013-01-15 07:09:01.000 | 655 |
+-----+-----+-----+-----+-----+
39 rows selected (0.591 seconds)
0: jdbc:phoenix:e3ecmrhdp11.mercury.corp,e3ec> █
```

3.3.3 Create a phoenix view for existing HBase Table.

```
Create View IM_ItemPrice for Testing
```

```
CREATE VIEW "data:basetable" ( pk VARCHAR PRIMARY KEY,  
"Info"."AITMark" VARCHAR,  
"Info"."Active" VARCHAR,  
"Info"."Buyer" VARCHAR,  
"Info"."ComboReservedQ4S" VARCHAR,  
"Info"."CompanyCode" VARCHAR,  
"Info"."CountryCode" VARCHAR,  
"Info"."CurrencyCode" VARCHAR,  
"Info"."DefaultShipVia" VARCHAR,  
"Info"."DefaultShippingCharge" VARCHAR,  
"Info"."DisableAutoActivateFlag" VARCHAR,  
"Info"."DiscountInstant" VARCHAR,  
"Info"."EggSaverMark" VARCHAR,  
"Info"."EggSaverMaxQty" VARCHAR,  
"Info"."EggSaverShippingCharge" VARCHAR,  
"Info"."EggSaverShippingType" VARCHAR,  
"Info"."GiftWrapItem" VARCHAR,  
"Info"."ItemNumber" VARCHAR,  
"Info"."LastEditDate" VARCHAR,  
"Info"."LastEditUser" VARCHAR,  
"Info"."LimitQuantity" VARCHAR,  
"Info"."MAPPrice" VARCHAR,  
"Info"."MAPPriceMark" VARCHAR,  
"Info"."ManufacturerWarrantyMark" VARCHAR,  
"Info"."MinShippingCharge" VARCHAR,  
"Info"."MinimumQuantity" VARCHAR,  
"Info"."PremierMark" VARCHAR,  
"Info"."PriceAutoAdjustMark" VARCHAR,  
"Info"."PriceHideMark" VARCHAR,  
"Info"."ProductManager" VARCHAR,  
"Info"."PromotionPrice1" VARCHAR,  
"Info"."PromotionPrice2" VARCHAR,  
"Info"."PromotionPrice3" VARCHAR,  
"Info"."PromotionPrice4" VARCHAR,  
"Info"."PromotionPrice5" VARCHAR,  
"Info"."PromotionQTY1" VARCHAR,  
"Info"."PromotionQTY2" VARCHAR,  
"Info"."PromotionQTY3" VARCHAR,  
"Info"."PromotionQTY4" VARCHAR,  
"Info"."PromotionQTY5" VARCHAR,  
"Info"."PromotionShipping1" VARCHAR,  
"Info"."PromotionShipping2" VARCHAR,
```

```
"Info"."PromotionShipping3" VARCHAR,  
"Info"."PromotionShipping4" VARCHAR,  
"Info"."PromotionShipping5" VARCHAR,  
"Info"."PurolatorCANShippingCharge" VARCHAR,  
"Info"."PurolatorCANShippingType" VARCHAR,  
"Info"."PurolatorUSAShippingCharge" VARCHAR,  
"Info"."PurolatorUSAShippingType" VARCHAR,  
"Info"."RepackMark" VARCHAR,  
"Info"."RestrictedItemMark" VARCHAR,  
"Info"."SameDayShippingCharge" VARCHAR,  
"Info"."SameDayShippingMark" VARCHAR,  
"Info"."SameDayShippingType" VARCHAR,  
"Info"."ShipByNewegg" VARCHAR,  
"Info"."ShippingPromotionGroupID" VARCHAR,  
"Info"."ShippingType" VARCHAR,  
"Info"."ShopRunnerMark" VARCHAR,  
"Info"."SignatureRequiredMark" VARCHAR,  
"Info"."UnitPrice" VARCHAR,  
"Info"."WarrantyCode" VARCHAR,  
"Info"."WebsiteBlockMark" VARCHAR  
);
```

3.3.4 Enlarge the time out seconds for HBase client

```
<property>
  <name>hbase.rpc.timeout</name>
  <value>600000000</value>
</property>
<property>
  <name>hbase.snapshot.enabled</name>
  <value>true</value>
</property>
<property>
  <name>hbase.snapshot.master.timeoutMillis</name>
  <value>600000000</value>
</property>
<property>
  <name>hbase.snapshot.region.timeout</name>
  <value>600000000</value>
</property>
<property>
  <name>hbase.snapshot.master.timeout.millis</name>
  <value>600000000</value>
</property>
<property>
  <name>hbase.security.authentication</name>
  <value>simple</value>
</property>
<property>
  <name>zookeeper.session.timeout</name>
  <value>600000000</value>
</property>
```

3.3.5 Testing on ItemPriceSetting data

Count IM_ItemPrice for 93097527 records in 315 seconds (5 minutes)

```
select count(1) from "data:basetable";
```

```
0: jdbc:phoenix:e3ecmrhdp11.mercury.corp,e3ec> select count(1) from "ecitem:IM_ItemPrice";
+-----+-----+
|          COUNT(1)          |
+-----+-----+
| 93097527                    |
+-----+-----+
1 row selected (315.634 seconds)
```

Filter entire table data on specific column in 6 minutes.

```
select count(1) from "data:basetable" where "Active"='False' AND "Buyer"='HOUS';
```

```
0: jdbc:phoenix:ecmrhdp11.mercury.corp,ec> select count(1) from "ecitem:IM_ItemPrice" where "Active"='False' AND "Buyer"='HOUS';
+-----+
| COUNT(1) |
+-----+
| 46787018 |
+-----+
1 row selected (382.031 seconds)
```

`select * from "data.basetable" where "Active"='False' limit 100;`

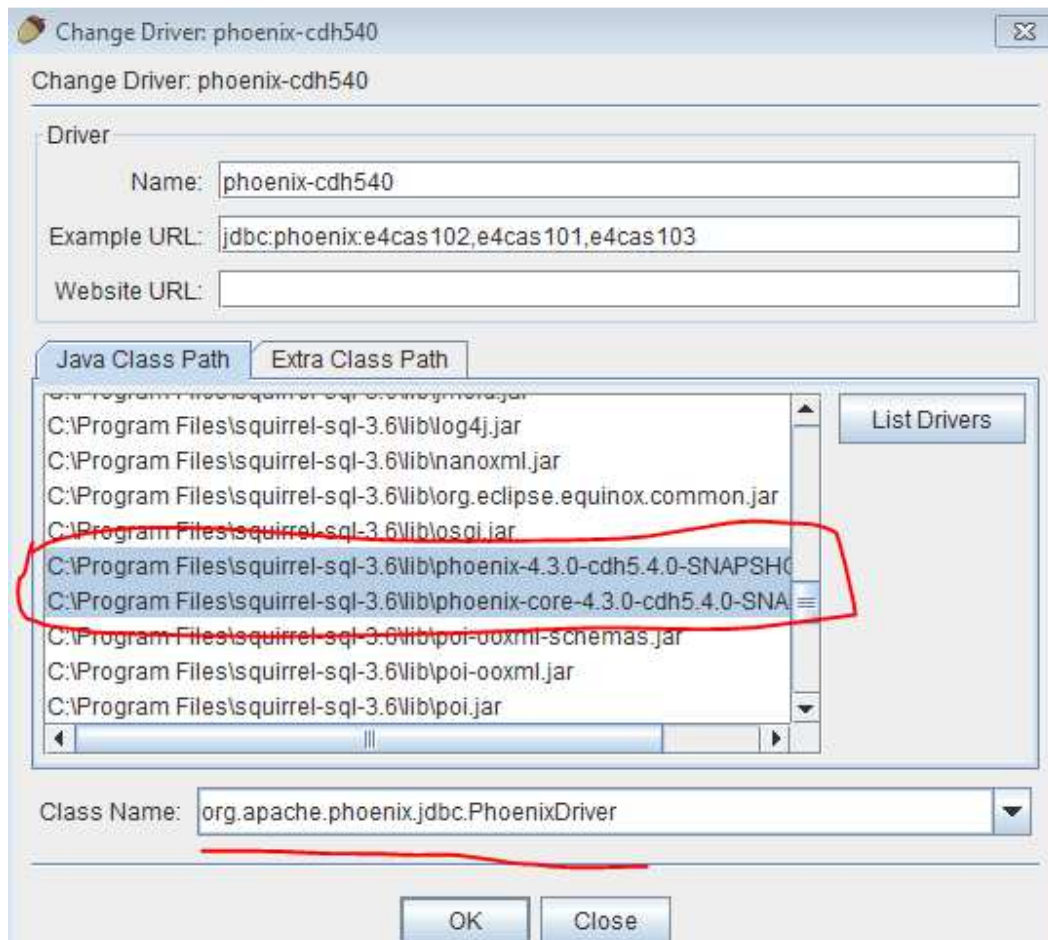
```
0: jdbc:phoenix:ecmrhdp11.mercury.corp,ec> select * from "ecitem:IM_ItemPrice" where "Active"='False' limit 100;
+-----+-----+-----+-----+
| PK | AITMark | Active | Buyer |
+-----+-----+-----+-----+
| 0000003cae0d0d170da4ec84bd7e06461f6000f8f1e307988f1f81fe3e86aa8b2 | False | False | HOUS |
| 0000006e7598538de7994a75e9cc0rec509ebf729115b496d754a8a3b791925a | False | False | HOUS |
| 000000c5a3962b662d25427858c766c2edea09d6d2723b460c64e5c466fa | False | False | CY36 |
| 0000016eadd647545b1fe1355bfc9830b70c1704dc4a2b72054353a0b830465d | False | False | HOUS |
| 00000189a0c5a03a68b7f94862d28c5cd31407efb2f3a84ddc547c92a836bf93 | False | False | HOUS |
| 0000019929bf28ac907e7f9dc5f5fe1bc7b227ff63fa61cf4ae0d2634a3e6 | False | False | HOUS |
| 0000019eb20074a1c0bdee53fa4b1d7a0d6c75d850d1f7cf0e23599c49d27b4e | False | False | hp56 |
| 000001aa45ee2d263523738f787192ffbae7c91a1f4017dd526789f029eb9b0b | False | False | KW35 |
| 000002660c0cc3c2861ec80ae37bef7b94baafeba3733de0bd3f89cc60c3107e | False | False | HOUS |
| 0000027f823a5967d2036eb7de34e6f83fae3d4cfd0640bf0bdc18adeda371a | False | False | HOUS |
| 000002ad3a1b4783224b39d3d7071fd10eb037852a5063b24100f9702b455d12d | False | False | cc01 |
| 0000033f0c395646d4a1736071778c9236ba3b0de0d1861a902ade9138070570 | False | False | HOUS |
| 0000038365c27ab6e49dcda49df8d2dcb16882f96a79fd388aec873c43528fb | False | False | HOUS |
| 000003eaeacac9596b90ef42c9ef3af2ec23ae7375cc019f9c1a001b2d14f2 | False | False | HOUS |
| 0000044b32704de054e121e0b9171fb69d8d1eb6d22561a83157ed0c88343fb | False | False | HOUS |
| 0000047fba2ba556c47775c519b40c5e95d793bf4f9bedcddf4282a6e09224c | False | False | HOUS |
| 00000490ea5a27151df9d20f5700176382da99e8fd0fd575f95e12468a76d2b3 | False | False | HOUS |
| 000004e21a3e8c619012812572c0d062ea2120950b9dc445eafb061599f85dd9 | False | False | HOUS |
| 000004fe984d8ea656fa9e57b8f28d21540396b0d1e4c1f365055cc5d7ec7659 | False | False | PY55 |
| 0000051cb440690f7a5199dc9ac7275a4a49fe470791202fa7f25a2d675ecaf6 | False | False | HOUS |
100 rows selected (4.501 seconds)
```

3.4 Use Squirrel Client

3.4.1 Add Phoenix Driver

Put phoenix-*.client.jar and phoenix-core-*.jar in <Squirrel Path>/lib

If you need to customize the hbase-site.xml /hdfs-site.xml / core-site/xml, you can also put them in <Squirrel Path>/lib path.



3.4.2 Add Alias for Phoenix Client:

Enter ZK Server List for URL.

Change Alias: px

Name: px

Driver: ✓ phoenix-cdh540 New

URL: jrp,e3ecmrhdp04.mercury.corp,e3ecmrhdp02.mercury.corp

User Name: tw79

Password:

☐ Auto logon ☐ Connect at Startup

Properties

Warning - Passwords are saved in clear text

OK Close Test

You can do a test for connection by click "Test" button.

Change Alias: px

Connect to: px

Alias: px

Driver: phoenix-cdh540

URL: jdbc:phoenix:e3ecmrhdp11.mercur...

User: tw79

Password:

Properties

Warning - Caps lock may interfere with passwords

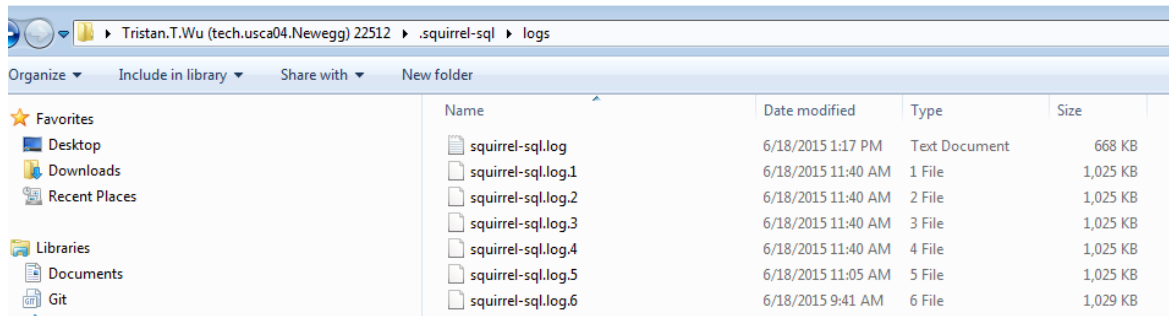
Connect Cancel

Connecting...

Connection successful

OK

You can check the logs from you <home folder>\squirrel-sql\logs.



3.4.3 Write a SQL to query phoenix table.

4 Summary

In my opinion, Apache Phoenix can only be used for ad-hoc queries in offline HBase.

4.1 Positive

- Apache Phoenix can support Ad-Hoc Query on HBase. The performance on 90 million PRD item data is 5-6 minutes.
- Apache Phoenix provides JDBC Driver, easy to use it.
- Apache Phoenix supports ANSI SQL.

4.2 Negative

- Apache Phoenix on Cludera needs to build against CDH5.4.0 by source code.

- Apache Phoenix is powered by HBase Coprocessor, which means scan / aggregate entire data on each region data. It cost lots of CPU. It is high risk for online HBase Cluster, only can use it in offline cluster.
- Not robust engine, and no rich use experience in community.