

Presto on CDH540 Testing Report

Tao Wu, 08/06/2015

1	Background	1
2	Environment.....	2
3	Testing Report	2
3.1	Check the cluster status	2
3.2	Simple count query	3
3.3	Simple distinct query.....	3
3.4	Condition query (filter).....	4
3.5	Select large result set	5
3.6	Complicated query with filter, like, group by, order by, count, top.	5
4	Server Pressure Monitor	6
4.1	Coordinator Status	7
4.2	Worker Status	7

1 Background

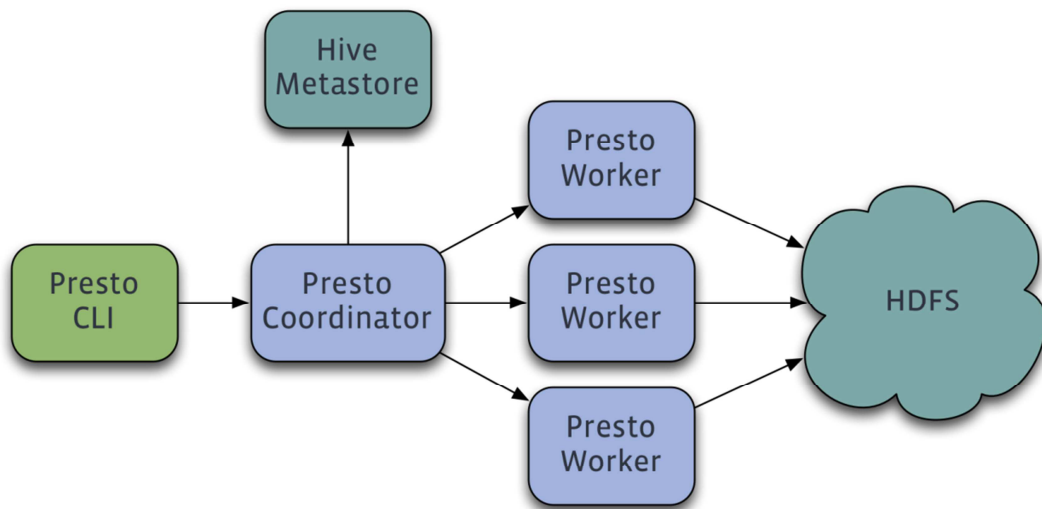
We have Hive (based on MapReduce V1 on CDH 4 version and MapReduce V2 on CDH5 version) to do SQL query and SQL job for offline data processing, it can process large-scale of data, but the time cost is minute-level.

Meanwhile, we still have lot of Storage Procedure in SQL Server.

As I mentioned before, we can make a long-term goal to transfer these kind of SQL SP job to big data. But the hive on MapReduce is no fast enough.

After one week testing and configuration optimization, I found the Presto is really good for SQL Query on Big Data. And I hope to give a demo show to EC team, to discuss what we can do on this fast SQL engine.

Presto is distributed SQL query engine open-sourced by Facebook, now Facebook has used it in PRD over two years, it is the upgrade version for hive in Facebook(hive is also developed and open-source by facebook). NetFlix, Aribnb, and TeraData are engaging in presto open-source community.



<https://prestodb.io/>

2 Environment

1. Centos 6.5 final
2. Oracle JDK1.8.0_51
3. Presto 0.112
4. Hadoop CDH540 (Hadoop-2.6.0-CDH, hive-1.1.0)
5. MySQL for Presto Meta Data
6. 16 node, 1 coordinator, 16 worker. 16 GB HeapSize for Presto Worker.

3 Testing Report

3.1 Check the cluster status

```
Select * from system.runtime.nodes;
```

```

presto:mars> select * from system.runtime.nodes;

```

node_id	http_uri	node_version	coordinator	active
de0ee2bd-4064-49a5-b033-04436822955b	http://172.16.31.155:8866	presto-main:0.112	false	true
31330ed9-bba2-4836-957e-bcde0cedf965	http://172.16.31.143:8866	presto-main:0.112	false	true
ba64ce7b-4850-4aff-a446-3a79cc603bd0	http://172.16.31.146:8866	presto-main:0.112	false	true
10a79203-dc12-4b10-862a-656b3939ed45	http://172.16.31.153:8866	presto-main:0.112	false	true
21b80c88-bea2-4f8c-b1c1-b5139db92a27	http://172.16.31.156:8866	presto-main:0.112	false	true
7df814c4-021b-4c60-9569-1753eb1f9d24	http://172.16.31.152:8866	presto-main:0.112	false	true
a2f1ed5f-dbd2-46d5-b699-82b7752cf3f7	http://172.16.31.144:8866	presto-main:0.112	false	true
ce1ccf2f-9cb6-4489-8ed2-cac880eae491	http://172.16.31.149:8866	presto-main:0.112	false	true
2c4ba944-cc59-407a-8240-3cffb795ecff	http://172.16.31.141:8866	presto-main:0.112	true	true
65e37afc-e439-4e44-b28b-92ff9248a1af	http://172.16.31.142:8866	presto-main:0.112	false	true
0d8ebb4b-20ea-459a-bc74-8d46194503d4	http://172.16.31.148:8866	presto-main:0.112	false	true
c336f28d-c02b-4366-a411-9bcfaf745553	http://172.16.31.147:8866	presto-main:0.112	false	true
dc76d898-8525-45d7-9987-8ebe9092fca0	http://172.16.31.145:8866	presto-main:0.112	false	true
07e4cdee-24ce-424d-939e-7bbf7ee7498b	http://172.16.31.154:8866	presto-main:0.112	false	true
8f0a52ba-6f58-4e95-99d5-480a4ce420d0	http://172.16.31.150:8866	presto-main:0.112	false	true
92251a97-390f-4449-a7b9-a453f4e4fdda	http://172.16.31.151:8866	presto-main:0.112	false	true

```

(16 rows)

Query 20150806_164207_00005_hrzhc, FINISHED, 2 nodes
Splits: 2 total, 2 done (100.00%)
0:01 [16 rows, 1.25KB] [19 rows/s, 1.51KB/s]

```

3.2 Simple count query

```
Select count(1) as rowcount from truesight_page_orc;
```

Done in 8 seconds, 831434942 rows.

```

presto:mars> select count(1) as rowcount from truesight_page_orc;
rowcount
-----
831434942
(1 row)

Query 20150806_164521_00007_hrzhc, FINISHED, 16 nodes
Splits: 3,872 total, 3,872 done (100.00%)
0:08 [831M rows, 290GB] [99.9M rows/s, 34.9GB/s]

```

3.3 Simple distinct query

```
Select distinct(dt) from truesight_page_orc;
```

Done in 20 seconds on 800 million records, 10 days data.

```
presto:mars> select distinct(dt) from truesight_page_orc;
dt
```

```
-----
```

```
20150729
20150728
20150803
20150727
20150731
20150804
20150801
20150802
20150730
20150726
(10 rows)
```

```
Query 20150806_164621_00008_hrzhc, FINISHED, 16 nodes
Splits: 3,880 total, 3,880 done (100.00%)
0:20 [831M rows, 290GB] [41.6M rows/s, 14.5GB/s]
```

3.4 Condition query (filter)

```
Select * from truesight_page_orc where dt='20150726' and hour='00' and
lower(cs_host) like 'secure%';
```

Done in 17 seconds on 800 million records.

```
presto:mars> select * from truesight_page_orc where dt='20150726' and hour='00' and lower(cs_host) like 'secure%';
```

x_location	x_record_type	x_page_id	x_session_id	x_server_id	x_page_name	cs_host	
EE	PAGE	54378100441	4484798514	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378124466	4484873264	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378116308	4484873830	172.16.20.146	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378112173	4484873315	172.16.20.147	orderhistory	secure.newegg.com	/NewMyAccount/OrderHistory.aspx
EE	PAGE	54378113669	4484873902	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378103734	4484098949	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378117063	4484873142	172.16.20.145	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378113355	4484873123	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378117688	4484873142	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378116632	4484874324	172.16.20.145	shoppingcart	secure.newegg.com	/Shopping/ShoppingCart.aspx
EE	PAGE	54378129519	4484873142	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378116071	4484382263	172.16.34.9	from	secure.newegg.com	/5505E12A-1DDD-B941-AA43-E6A8DAFF
EE	PAGE	54379804149	4485002288	172.16.20.147	lastpurchaseddate	secure.neweggbusiness.com	/new/ajax/lastpurchaseddate
EE	PAGE	54379803651	4484382263	172.16.34.9	from	secure.newegg.com	/5505E12A-1DDD-B941-AA43-E6A8DAFF
EE	PAGE	54379802724	4485000245	172.16.34.7	checkoutstep3	secure.newegg.com	/GlobalShopping/CheckoutStep3.aspx
EE	PAGE	54379800273	4484944945	172.16.34.7	list	secure.ows.newegg.com	/customers/order.egg/list
EE	PAGE	54379804938	4484847364	172.16.34.5	shoppingcart	secure.newegg.com	/Shopping/ShoppingCart.aspx
EE	PAGE	54379874666	4484914362	172.16.20.134	38247312	secure.ows.newegg.com	/Premier.egg/IsPremierMembership
EE	PAGE	54379806078	4484113232	172.16.20.134	wishlist	secure.newegg.ca	/WebResource/Scripts/CAN/SSL/Wish
EE	PAGE	54379805421	4484382263	172.16.34.9	from	secure.newegg.com	/5505E12A-1DDD-B941-AA43-E6A8DAFF
EE	PAGE	54379077923	4484913651	172.16.20.145	40945812	secure.ows.newegg.com	/Premier.egg/IsPremierMembership
EE	PAGE	54379805980	4484947560	172.16.20.134	apple-touch-icon-precomposed	secure.newegg.com	/apple-touch-icon-precomposed.png
EE	PAGE	54378100585	4484768038	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378113882	4484871797	172.16.34.6	securecheckout	secure.m.newegg.com	/Account/Login
EE	PAGE	54378117228	4484873222	172.16.20.145	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378117344	4484873538	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378116528	4484874017	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378104770	4484871935	172.16.20.146	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378120793	4484873538	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378121011	4484873142	172.16.20.145	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378117115	4484873902	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378120353	4484873123	172.16.20.134	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378105218	4484872126	172.16.34.9	37511031	secure.ows.newegg.com	/Premier.egg/IsPremierMembership
EE	PAGE	54378125344	4484873123	172.16.20.146	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378108597	4484872759	172.16.20.147	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378109185	4484873915	172.16.20.145	shoppingcart	secure.newegg.com	/Shopping/ShoppingCart.aspx
EE	PAGE	54378110714	4484382263	172.16.34.9	from	secure.newegg.com	/5505E12A-1DDD-B941-AA43-E6A8DAFF
EE	PAGE	54378116909	4484873498	172.16.20.145	accountlogin	secure.newegg.com	/NewMyAccount/AccountLogin.aspx
EE	PAGE	54378092268	4484869580	172.16.20.134	shoppingcart	secure.newegg.com	/Shopping/ShoppingCart.aspx
EE	PAGE	54378111269	4484818531	172.16.34.7		secure.ows.newegg.com	/wishlist.egg/detail/

```

EE | PAGE | 54378133801 | 4484873498 | 172.16.20.134 | accountlogin | secure.newegg.com | /NewMyAccount/AccountLogin.aspx
EE | PAGE | 54378112260 | 4484871576 | 172.16.20.145 | accountlogin | secure.newegg.com | /NewMyAccount/AccountLogin.aspx
EE | PAGE | 54378117479 | 4484874017 | 172.16.20.134 | accountlogin | secure.newegg.com | /NewMyAccount/AccountLogin.aspx
(query aborted by user)

Query 20150806 165234 00010 hrzhc, RUNNING, 16 nodes
Splits: 3,872 Total, 1,120 Done (28.93%)
0:17 [420M rows, 165GB] [24.5M rows/s, 9.64GB/s]

```

3.5 Select large result set

```
Select * from truesight_page_orc where dt='20150726' and hour='00' and
lower(cs_host) like 'secure%';
```

Done in 25 seconds on 800 million records.

```

presto:marco> select * from truesight_page_orc where dt='20150726' and hour='00' and lower(cs_host) like 'secure%';
x_location | x_record_type | x_page_id | x_session_id | x_server_id | x_page_name | cs_host |
-----
EE | PAGE | 54380711325 | 4484382263 | 172.16.34.9 | from | secure.newegg.com | /5505E12A-1DD0-B941-AA43-E6A8DAF
EE | PAGE | 54380694086 | 4484254168 | 172.16.34.9 | 40101472 | secure.ows.newegg.com | /Premier.egg/IsPremierMembership
EE | PAGE | 54380706500 | 4484950449 | 172.16.34.5 | 2512026 | secure.ows.newegg.com | /Premier.egg/IsPremierMembership
EE | PAGE | 54380714111 | 4484382263 | 172.16.34.9 | from | secure.newegg.com | /5505E12A-1DD0-B941-AA43-E6A8DAF
EE | PAGE | 54380710694 | 4484461175 | 172.16.34.1 | 36065525 | secure.ows.newegg.com | /Premier.egg/IsPremierMembership
EE | PAGE | 54380714399 | 4485063727 | 172.16.34.5 | addtocart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380693291 | 4485052483 | 172.16.20.134 | checkoutstep2 | secure.newegg.com | /Guest/CheckoutStep3.aspx
EE | PAGE | 54380711300 | 4484888246 | 172.16.34.6 | accountlogin | secure.neweggbusiness.com | /New/MyAccount/AccountLogin.aspx
EE | PAGE | 54380711914 | 4485067469 | 172.16.20.134 | accountlogin | secure.newegg.com | /NewMyAccount/AccountLogin.aspx
EE | PAGE | 54380714955 | 4484382263 | 172.16.34.9 | from | secure.newegg.com | /5505E12A-1DD0-B941-AA43-E6A8DAF
EE | PAGE | 54380717433 | 4485067528 | 172.16.20.134 | accountlogin | secure.newegg.com | /NewMyAccount/AccountLogin.aspx
EE | PAGE | 54380708533 | 4485063527 | 172.16.20.145 | 49677850 | secure.ows.newegg.com | /Premier.egg/IsPremierMembership
EE | PAGE | 54380690159 | 4485062257 | 172.16.34.7 | shoppingcart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380688635 | 4485065550 | 172.16.20.134 | shoppingcart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380689603 | 4484113232 | 172.16.20.134 | %5Cetc%5Chosts | secure.newegg.ca | /WebResource/Themes/2005/CSS/Sec
EE | PAGE | 54380689225 | 4485066336 | 172.16.20.145 | addtocart | secure.newegg.ca | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380690815 | 4485066305 | 172.16.34.6 | iframe3rdpartycodeTotal | secure.newegg.com | /Common/ThirdParty/IFrame3rdPart
EE | PAGE | 54380688076 | 4485052145 | 172.16.20.145 | createpassword | secure.newegg.ca | /NewMyAccount/MessagePage.aspx
EE | PAGE | 54380688110 | 4485065132 | 172.16.20.146 | iframe3rdpartycodeTotal | secure.newegg.com | /Common/ThirdParty/IFrame3rdPart
EE | PAGE | 54380690521 | 4485066223 | 172.16.34.7 | minicartupdate | secure.newegg.com | /Common/Ajax/MiniCartUpdate.aspx
EE | PAGE | 54380688541 | 4485061319 | 172.16.20.145 | paypalreturn | secure.newegg.com | /Guest/CheckoutStep3.aspx
EE | PAGE | 54380688758 | 4485063374 | 172.16.20.147 | temporarywishlist | secure.newegg.com | /Wishlist/TemporaryWishlist.aspx
EE | PAGE | 54380689396 | 4485066336 | 172.16.20.145 | shoppingcart | secure.newegg.ca | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380689394 | 4485065972 | 172.16.20.147 | iframe3rdpartycodeTotal | secure.newegg.com | /Common/ThirdParty/IFrame3rdPart
EE | PAGE | 54380688091 | 4485062237 | 172.16.20.134 | wishcartplus | secure.newegg.com | /Wishlist/TemporaryWishlist.aspx
EE | PAGE | 54380691533 | 4485066161 | 172.16.34.9 | checkstatebyp | secure.newegg.com | /Common/Ajax/CheckStateByIP.aspx
EE | PAGE | 54380689175 | 4485054927 | 172.16.20.146 | shoppingcart | secure.newegg.com | /Shopping/ShoppingLogin.aspx
EE | PAGE | 54380689316 | 4485064798 | 172.16.20.147 | checkoutstep3 | secure.newegg.com | /GlobalShopping/CheckoutStep3.as
EE | PAGE | 54380691202 | 4485064221 | 172.16.34.9 | accountlogin | secure.newegg.com | /NewMyAccount/AccountLogin.aspx
EE | PAGE | 54380685670 | 4485063751 | 172.16.20.145 | addtocart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380688381 | 4485065972 | 172.16.20.147 | addtocart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380688115 | 4485063751 | 172.16.20.145 | iframe3rdpartycodeTotal | secure.newegg.com | /Common/ThirdParty/IFrame3rdPart
EE | PAGE | 54380689597 | 4485018664 | 172.16.20.145 | orderhistory | secure.newegg.com | /NewMyAccount/OrderHistory.aspx
EE | PAGE | 54380690766 | 4485066293 | 172.16.34.9 | shoppingcart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380688439 | 4485065550 | 172.16.20.134 | shoppingcart | secure.newegg.com | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380688324 | 4484795152 | 172.16.20.145 | secure.ows.newegg.com | /wishlist.egg/detail/
EE | PAGE | 54380689083 | 4485062453 | 172.16.20.134 | responseorderfrombank | secure.newegg.com | /Shopping/CheckoutStep3.aspx
EE | PAGE | 54380688127 | 4485066180 | 172.16.20.145 | addtocart | secure.newegg.ca | /Shopping/ShoppingCart.aspx
EE | PAGE | 54380689562 | 4485062774 | 172.16.20.134 | messages | secure.newegg.com | /NewMyAccount/Messages.aspx
EE | PAGE | 54380691362 | 4484939415 | 172.16.34.7 | shoppingcartIsDirectCheckout | secure.ows.newegg.com | /api/ShoppingCartIsDirectCheckou
EE | PAGE | 54380690702 | 4484924703 | 172.16.34.6 | shoppingcart | secure.newegg.com | /Shopping/ShoppingCart.aspx

EE | PAGE | 54379544556 | 4484980610 | 172.16.34.6 | iframe3rdpartycodeTotal | secure.newegg.com | /Common/ThirdParty/IFrame3rdParty
EE | PAGE | 54379544386 | 4484978030 | 172.16.20.134 | 21447005 | secure.ows.newegg.com | /Premier.egg/IsPremierMembership/
EE | PAGE | 54379546180 | 4484979425 | 172.16.34.5 | subscriptions | secure.newegg.ca | /NewMyAccount/MessagePage.aspx
EE | PAGE | 54380245599 | 4483945871 | 172.16.34.9 | accountlogin | secure.neweggbusiness.com | /new/newmyaccount/accountlogin.as
(query aborted by user)

Query 20150806 165536 00014 hrzhc, RUNNING, 16 nodes
Splits: 3,072 Total, 595 Done (15.37%)
0:25 [317M rows, 129GB] [12.8M rows/s, 5.19GB/s]

```

3.6 Complicated query with filter, like, group by, order by, count, top.

```

select
c_ip,
cs_uri_stem,
cs_host,
count(cs_uri_stem) ct,
cs_method,
sc_status,

```

```
count( distinct x_session_id) sessioncount
from oos.traffic_data_orc
where lower(cs_host) like 'secure%'
and c_ip<>' ' and c_ip <> '198.100.137.34'
group by c_ip, cs_uri_stem,cs_host,cs_method,sc_status
having count(cs_uri_stem)>500
limit 100;
```

Done in 1 minute 10 seconds on 800 million records.

```
presto:mars> select c_ip, cs_uri_stem, cs_host, count(cs_uri_stem) ct, cs_method, sc_status, count( distinct x_session_id) sessioncount from mars.truesight_page_orc
ere lower(cs_host) like 'secure%' and c_ip<>' ' and c_ip <> '204.14.213.149' group by c_ip, cs_uri_stem,cs_host,cs_method,sc_status having count(cs_uri_stem)>500 orde
by ct desc limit 100;
```

c_ip	cs_uri_stem	cs_host	ct	cs_method	sc_status	ses
184.173.129.66	/Capcha/Index	secure.m.newegg.com	98818	GET	200	
74.86.197.38	/Shopping/ShoppingCart.aspx	secure.newegg.com	60287	GET	200	
50.17.119.107	/new/newmyaccount/b2baccountregistration.aspx	secure.neweggbusiness.com	28437	GET	302	
50.17.119.107	/new/newmyaccount/accountlogin.aspx	secure.neweggbusiness.com	28434	GET	302	
74.62.251.236	/5505E12A-10DD-B941-AA43-E6A8DAFF4410/88F27D2C-8ABD-D14A-89F8-E414AE28F19A/from/	secure.newegg.com	22288	GET	404	
54.237.246.232	/new/newmyaccount/accountlogin.aspx	secure.neweggbusiness.com	21300	GET	302	
54.237.246.232	/new/newmyaccount/b2baccountregistration.aspx	secure.neweggbusiness.com	21292	GET	302	
206.190.149.45	/Shopping/ShoppingCart.aspx	secure.newegg.ca	20107	GET	200	
54.209.60.63	/NewMyAccount/AccountLogin.aspx	secure.newegg.com	19051	GET	200	
73.231.112.43	/A5CB5E92-38C3-3340-8F97-D8FB99C6A833/467BB893-5028-FF4C-9D03-C214AC447B20/from/	secure.newegg.com	17681	GET	404	
54.83.105.242	/new/newmyaccount/b2baccountregistration.aspx	secure.neweggbusiness.com	15948	GET	302	
54.83.105.242	/new/newmyaccount/accountlogin.aspx	secure.neweggbusiness.com	15943	GET	302	
54.83.206.213	/Shopping/AddToCart.aspx	secure.newegg.com	14629	GET	302	
50.17.119.107	/new/feedback/rateproduct.aspx	secure.neweggbusiness.com	12512	GET	302	
54.209.60.63	/Shopping/ShoppingCart.aspx	secure.newegg.com	12394	GET	200	
24.236.250.254	/New/404.aspx	secure.neweggbusiness.com	12100	GET	302	
107.145.126.131	/CB13D742-B070-844E-9B19-97E8D7803993/945B30F5-3345-6446-9275-3889E67860D3/from	secure.newegg.com	12052	POST	302	
54.83.206.213	/Shopping/ShoppingCart.aspx	secure.newegg.com	11642	GET	302	
71.190.132.28	/New/NewMyAccount/AccountLogin.aspx	secure.neweggbusiness.com	11206	GET	200	
216.52.208.149	/Common/Ajax/TrackOrderHandler.aspx	secure.newegg.com	10889	GET	200	
71.21.34.177	/GiftCard/BalanceCheck.aspx	secure.newegg.com	9954	POST	200	
202.133.226.43	/ShoppingLogin.egg	secure.ows.newegg.com	9629	POST	200	
54.237.246.232	/new/feedback/rateproduct.aspx	secure.neweggbusiness.com	9475	GET	302	
54.241.198.78	/Shopping/AddToCart.aspx	secure.newegg.com	9105	HEAD	302	
50.18.94.121	/Shopping/AddToCart.aspx	secure.newegg.com	8800	HEAD	302	
54.151.42.39	/Shopping/AddToCart.aspx	secure.newegg.com	8142	HEAD	302	
174.115.118.120	/WebResource/Themes/2005/CSS/CAN/font-awesome.v1.s.10706.2.css	secure.newegg.ca	7490	HEAD	200	
204.89.152.90	/new/ajax/lastpurchaseddate	secure.neweggbusiness.com	6795	GET	200	
162.17.27.238	/Guest/OrderStatus.aspx	secure.newegg.com	6638	GET	200	
172.56.12.76	/ShoppingCart	secure.m.newegg.com	6276	GET	NULL	
54.83.105.242	/new/feedback/rateproduct.aspx	secure.neweggbusiness.com	6222	GET	302	
162.17.27.238	/Guest/OrderLogin.aspx	secure.newegg.com	5951	POST	302	
54.83.105.242	/new/newmyaccount/accountlogin.aspx	secure.neweggbusiness.com	5812	GET	200	
54.83.105.242	/new/newmyaccount/b2baccountregistration.aspx	secure.neweggbusiness.com	5812	GET	200	
54.83.206.213	/Shopping/ShoppingCart.aspx	secure.newegg.com	5770	GET	200	
204.89.152.91	/Common/ThirdParty/IFrame3rdPartyCodeTotal.aspx	secure.newegg.com	5669	GET	200	
107.145.126.131	/CB13D742-B070-844E-9B19-97E8D7803993/945B30F5-3345-6446-9275-3889E67860D3/from/	secure.newegg.com	5278	GET	404	
54.237.246.232	/new/newmyaccount/accountlogin.aspx	secure.neweggbusiness.com	4850	GET	200	
54.237.246.232	/new/newmyaccount/b2baccountregistration.aspx	secure.neweggbusiness.com	4845	GET	200	
24.189.141.202	/Shopping/ShoppingCart.aspx	secure.newegg.com	2264	GET	200	
198.205.118.69	/Shopping/ShoppingCart.aspx	secure.newegg.com	2245	GET	200	
24.236.250.254	/new/shopping/shoppinglogin.aspx	secure.neweggbusiness.com	2226	GET	200	
24.236.250.254	/new/newmyaccount/accountlogin.aspx	secure.neweggbusiness.com	2223	GET	200	
72.52.77.18	/Shopping/ShoppingCart.aspx	secure.newegg.com	2174	GET	200	
54.237.246.232	/New/NewMyAccount/AccountLogin.aspx	secure.neweggbusiness.com	2165	GET	200	
24.188.80.93	/Shopping/ShoppingCart.aspx	secure.newegg.com	2149	GET	200	
204.89.152.90	/new/ajax/email	secure.neweggbusiness.com	2134	GET	200	
208.94.246.251	/Shopping/ShoppingCart.aspx	secure.newegg.com	2131	GET	200	

(100 rows)

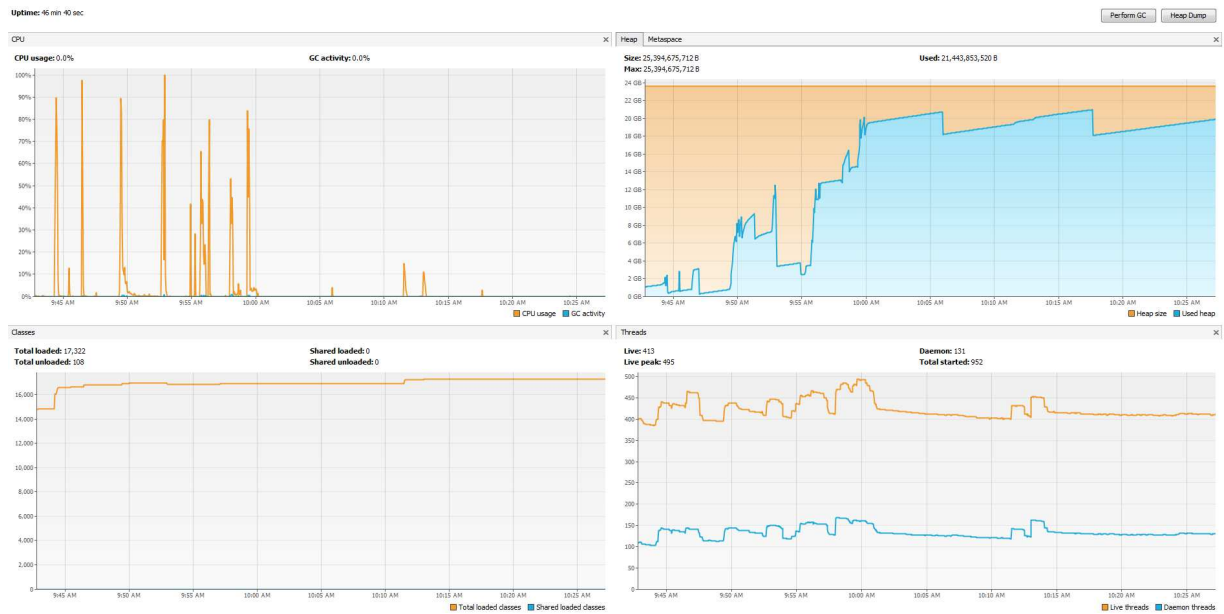
Query 20150806 165754 00016 hrzhc, FINISHED, 16 nodes
Splits: 3,888 total, 3,888 done (100.00%)
1:10 [831M rows, 290GB] [11.8M rows/s, 4.13GB/s]

4 Server Pressure Monitor

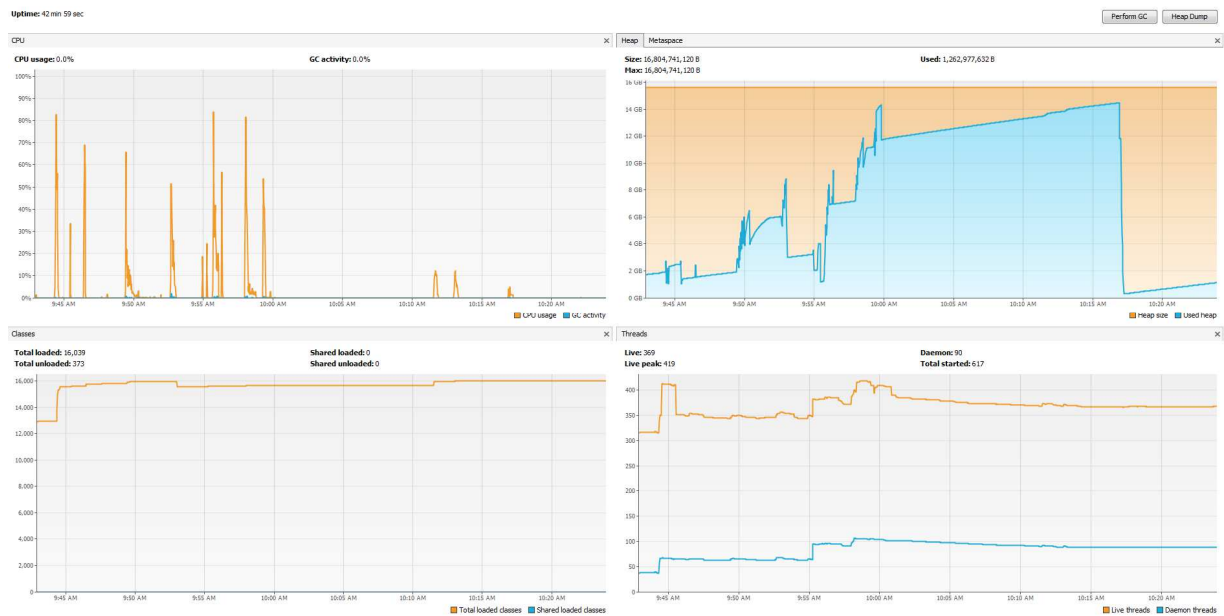
CPU will get peak when heavy query, but just a few seconds.

Memory will increase slowly, and GC can handle it.

4.1 Coordinator Status



4.2 Worker Status



4. ORC File vs CSV File

CSV File costs more CPU, and job may fail. ORC file can provide better and more stable performance.

```
select  
c_ip,  
cs_uri_stem,  
cs_host,
```

```

count(cs_uri_stem) ct,
cs_method,
sc_status,
count( distinct x_session_id) sessioncount
from oos.traffic_data
where dt like '2015072%' and lower(cs_host) like 'secure%'
and c_ip<>' ' and c_ip <> '198.100.137.34'
group by c_ip, cs_uri_stem,cs_host,cs_method,sc_status
having count(cs_uri_stem)>500
limit 1000;

```

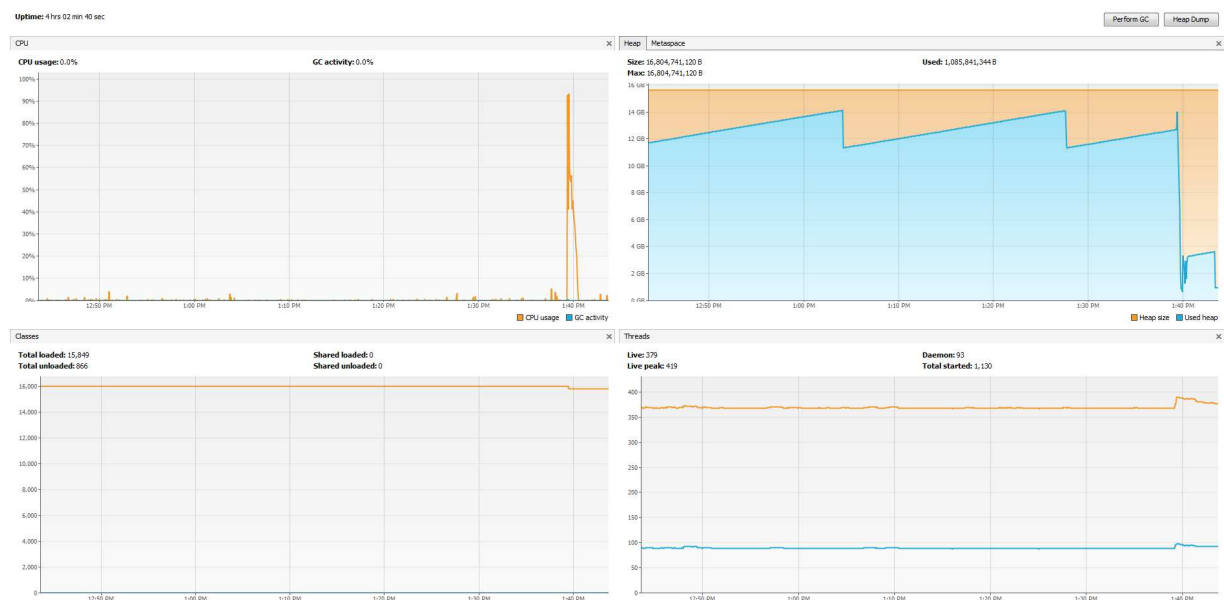
```

[tw79@e3ecmrhdp11 prc ip, cs_uri_stem, cs_host, count(cs_uri_stem) ct, cs_method, sc_status, count( distinct x_session_id) sessioncount from mars.truesight_page where
dt like '2015072%' and lower(cs_host) like 'secure%' and c_ip<>' ' and c_ip <> '204.14.213.149' group by c_ip, cs_uri_stem,cs_host,cs_method,sc_status having count(cs_u
ri_stem)>500 order by ct desc limit 100;

Query 20150806 203911 00023 hrzhc, FAILED, 16 nodes
Splits: 2,891 total, 1,900 done (65.72%)
1:10 [617M rows, 320GB] [8.83M rows/s, 4.57GB/s]

Query 20150806 203911 00023 hrzhc failed: Encountered too many errors talking to a worker node. The node may have crashed or be under too much load. This is probably a
transient issue, so please retry your query in a few minutes. (http://172.16.31.141:8866/v1/task/20150806_203911_00023_hrzhc.2.3/results/20150806_203911_00023_hrzhc.1
.1/0 - requests failed for 60.00s)

```



The coordinator may hang, and cannot support more queries.